

Distributed Information Systems

Fall Semester – 2023

CS-423

Time and Place

Lecture: Thursday 10:15 – 12:00, CM3

<https://epfl.zoom.us/j/66237387610>

Exercise: Thursday 12:15-13:00, CM3

Karl Aberer

Distributed Information Systems Laboratory

Goals of the Course

Understand what is a "**Distributed Information System**"?

- e.g. Web Search Engines, Online Social Networks, etc.

Know which are **key tasks** relevant for DIS?

- e.g. retrieval, mining, recommending, information extraction, data integration etc.

Master **common methods** used to solve these problems

- e.g. vector space model, graph mining, word embeddings etc.

Pre-existing knowledge not required

Knowledge in databases and machine learning helpful

Disclaimer on GPT (LLMs)

Many of these task can now be done with LLMs!

Why care about other methods?

- ▶ Understand the principles
- ▶ Alternative methods remain useful

Related Courses

- Applied Data Analysis
- Introduction to database systems
- Database systems
- Introduction to machine learning
- Machine learning
- Introduction to natural language processing
- Internet analytics
- ...

Which master's program are you from?

1. Computer Science
2. Communications
3. Data Science
4. Cybersecurity
5. Digital Humanities
6. Life Science
7. Electrical Engineering
8. Environmental Science
9. Others

The Course - Lecture

Standard online ex cathedra lecture

- Lecture streamed via Zoom
 - <https://epfl.zoom.us/j/66462767931>
 - Zoom QA tool to ask questions
 - Will be answered privately by assistants, or by the lecturer, depending on the questions
- Zoom Quizzes (anonymous)
- Zoom Chat to collect feedbacks

Video recordings

- Check on Moodle

Are you planning to join the lecture live or virtually?

1. I join live today, and plan to continue live
2. I join live today, but plan to join virtually
3. I join virtually today, and plan to continue virtually
4. I join virtually today, but plan to join live

Materials

Web platform: Moodle

- General announcements will be published on Moodle
- Course notes and project-related information will be published on the Web:
<https://lsir.github.io/DIS/>
- Exercises and exam questions from previous years will be made available as well

Projects - Key element of the course

Projects

1. Information retrieval
2. Recommender systems
3. Named entity recognition
4. Information extraction

Done in groups of 2 or 3

More details in exercise session

Project Evaluation

Results - Metrics [Comparison with baselines]

Code

- Working code
- Code quality and documentation

2-page Report [Moodle submission]

- Originality of approach
- Interpretation of results
- Presentation

Exercise Platform

Ed Forum to ask questions offline:

<https://edstem.org/eu/courses/831/discussion/>

Both among students and with assistants

Grading

Projects: 50%

Final Exam: 50%

- Program problem similar to the projects
- will assume you attended the lecture
- will assume you did the projects
- examples from earlier years (exercises, exams)
provided for preparation

Exam Support

Your computer will be admitted to the exam

- You will have Internet access
- But: communication not allowed (messaging, social platform etc.)
- You can use your notes (paper or electronically, all lecture materials)

Schedule

Week	Date	Area	Topic	Project
1	Thursday, 21 September 2023	Introduction		
2	Thursday, 28 September 2023	Information Retrieval	Basic Information Retrieval	Information Retrieval
3	Thursday, 5 October 2023	Embedding techniques		
4	Thursday, 12 October 2023	Embedding techniques		
5	Thursday, 19 October 2023	Recommender Systems	Recommender Systems	Recommender Systems
6	Thursday, 26 October 2023	Information Extraction	Named Entity Recognition	
7	Thursday, 2 November 2023	Document Classification		
8	Thursday, 9 November 2023	Knowledge Representation		Named Entity Extraction
9	Thursday, 16 November 2023	Information Extraction		
10	Thursday, 23 November 2023	Taxonomy induction		
11	Thursday, 30 November 2023	Graph Analytics	Link-based Ranking	Information Extraction
12	Thursday, 7 December 2023	Knowledge Inferences		
13	Thursday, 14 December 2023	Data Indexing and mining	Indexing for Information retrieval	
14	Thursday, 21 December 2023	Association Rule Mining		

Lecturer



Organizational Info

Moodle

- <http://moodle.epfl.ch/course/view.php?id=4051>

Lecturers

- Prof. Karl Aberer karl.aberer@epfl.ch BC 108

Assistants

- Romanou Angelika angelika.romanou@epfl.ch
- Negar Foroutan negar.foroutan@epfl.ch
- Mohammadreza Banaei mohammadreza.banaei@epfl.ch
- William Cappelletti william.cappelletti@epfl.ch


References

Parts of the course are based on the following text books


- Ricardo Baeza-Yates, Berthier Ribeiro-Neto, Modern Information Retrieval (Acm Press Series), Addison Wesley, 1999.
- Jiawei Han, Data Mining: concepts and techniques, Morgan Kaufman, 2000.
- Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze, Introduction to Information Retrieval, Cambridge University Press. 2008.
- J Leskovec, A Rajaraman, JD Ullman, Mining of Massive Datasets, 2014.

Further references to the literature will be given during the lecture

Free books



mmds.org



<http://nlp.stanford.edu/IR-book/>

Exam Date

Part 1 Information Retrieval

Part 1: Information Retrieval

1.1 Introduction to Information Retrieval

1.2 Basic Information Retrieval

1.2.1 Text-Based Information Retrieval

1.2.2 Boolean Retrieval

1.2.3 Vector Space Retrieval

1.2.4 Probabilistic Information Retrieval

1.2.5 Evaluating Information Retrieval

1.2.6 Query Expansion

1.2.6.1 User Relevance Feedback

1.2.6.2 Global Query Expansion

1.3 Embedding Techniques

1.3.1 Latent Semantic Indexing

1.3.2 Latent Dirichlet Allocation

1.3.3 Word Embeddings – skipgram, CBOW

1.3.4 Fasttext

1.3.5 Glove

Part 2: Recommender Systems

Part 2: Recommender Systems

- | | |
|-----|-----------------------------------|
| 2.1 | Collaborative Filtering |
| 2.2 | Content-based Recommendation |
| 2.3 | Matrix Factorization |
| 2.4 | SLIM, Sparse Linear Methods |
| 2.5 | Evaluation of Recommender Systems |

Part 3: Information Extraction

Part 3: Information Extraction

3.1 Named Entity Recognition

3.1.1 Keyphrase extraction

3.1.2 Named entity recognition (NER)

3.1.3 Entity Disambiguation

3.2 Document Classification

3.2.1 kNN

3.2.2 Naïve Bayes Classifier

3.2.3 Classification Using Word Embeddings

3.2.4 Transformer Models

3.3 Knowledge Representation

3.3.1 Knowledge Representation

3.3.2 Semi-structured data

3.3.3 The Semantic Web

3.3.4 RDF - Resource Description Framework

3.3.5 Semantic Web Resources

3.4 Information Extraction

3.4.1 Information extraction (IE)

3.4.1.1 Hand-written patterns

3.4.1.2 Supervised machine learning

3.4.1.3 Bootstrapping

3.4.1.4 Distant supervision

3.4.1.5 Matrix Factorization

3.4.2 Taxonomy Induction

Part 4: Graph Analytics

Part 4: Graph analytics

4.1 Link-Based Ranking

4.1.1 PageRank

4.1.2 Hyperlink-Induced Topic Search (HITS)

4.2 Mining Social Graphs

4.2.1 Louvain Modularity Algorithm

4.2.2 Girvan-Newman Algorithm

4.2 Knowledge Inference

4.2.1 Label Propagation

4.2.2 Link Prediction

Part 5: Data Indexing and Mining

Part 5: Data Indexing and Mining

5.1 Indexing for Information Retrieval

5.1.1 Inverted Index

5.1.2 Web-scale Indexing: Map-Reduce

5.1.3 Link Indexing

5.1.4 Distributed Retrieval

5.1.4.1 Fagin's algorithm

5.1.4.2 Threshold algorithm

5.2 Introduction to Data Mining

5.2.1 Association Rule Mining

5.2.1.1 Association Rules

5.2.1.2 Scoring Function

5.2.1.3 Apriori Algorithm

5.2.1.4 FP Growth